

dqn-cartpole

August 22, 2019

```
[1]: # https://github.com/higgsfield/RL-Adventure/blob/master/1.dqn.ipynb
# DQN without a frozen target network

[2]: %matplotlib inline
# %load_ext memory_profiler
# %load_ext line_profiler
# %load_ext heat
# %load_ext snakeviz

[3]: import yaml
import datetime

# from IPython.display import clear_output
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()

[4]: from torch.utils.tensorboard import SummaryWriter
# %reload_ext tensorboard
# %tensorboard --port=9706 --logdir ./runs
from torchsummary import summary

[5]: experiment_no = 'base_config'
# FROM CONFIG FILE
config_path = './' + experiment_no + '.yaml' # sys.argv[2]
config = yaml.safe_load(open(config_path, 'r'))

seed_value = 324267 # sys.argv[1]

# # Writer will output to ./runs/ directory by default
writer_dir = './runs/' + config['MODEL_NAME'] + '_' + experiment_no + '_' + \
    str(seed_value) + '_' + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
writer = SummaryWriter(writer_dir)
print("EXPERIMENT: ", experiment_no, "\tSEED: ", seed_value, "\twriter_dir: ", \
    writer_dir)
```

```
EXPERIMENT: base_config      SEED: 324267   writer_dir:
./runs/D2QN_base_config_324267_20190822-181003
```

```
[6]: import math
import os
import random
import numpy as np
import tensorflow as tf

import torch
import torch.nn as nn
import torch.optim as optim
import torch.autograd as autograd
import torch.nn.functional as F

[7]: os.environ['PYTHONHASHSEED']=str(seed_value)
random.seed(seed_value)
np.random.seed(seed_value)
# tf.random.set_seed(seed_value)
torch.manual_seed(seed_value)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = False
```

```

AttributeError                                Traceback (most recent call
last)

<ipython-input-7-c99862a89641> in <module>
      2 random.seed(seed_value)
      3 np.random.seed(seed_value)
----> 4 tf.random.set_seed(seed_value)
      5 torch.manual_seed(seed_value)
      6 torch.backends.cudnn.deterministic = True

~/anaconda3/envs/torchflow/lib/python3.6/site-packages/tensorflow/python/
util/deprecation_wrapper.py in __getattr__(self, name)
    104     if name.startswith('_dw_'):
    105         raise AttributeError('Accessing local variables before they
are created.')
--> 106     attr = getattr(self._dw_wrapped_module, name)
    107     if (self._dw_warning_count < _PER_MODULE_WARNING_LIMIT and
    108         name not in self._dw_deprecated_printed):

AttributeError: module 'tensorflow._api.v1.random' has no attribute
'set_seed'
```

```
[ ]: import gym
# CartPole-v0 Environment
env_id = "CartPole-v0"
env = gym.make(env_id)
env.seed(seed_value);

[ ]: USE_GPU = config['USE_GPU']

# Use CUDA
USE_CUDA = torch.cuda.is_available() and USE_GPU

if USE_CUDA:
    torch.cuda.manual_seed(seed_value)
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

[ ]: # REPLAY BUFFER

from collections import deque

class ReplayBuffer(object):
    def __init__(self, capacity):
        self.buffer = deque(maxlen=capacity)

    def push(self, state, action, reward, next_state, done):
        state = np.expand_dims(state, 0)
        next_state = np.expand_dims(next_state, 0)

        self.buffer.append((state, action, reward, next_state, done))

    def sample(self, batch_size):
        state, action, reward, next_state, done = zip(*random.sample(self.
→buffer, batch_size))
        return np.concatenate(state), action, reward, np.
→concatenate(next_state), done

    def __len__(self):
        return len(self.buffer)

[ ]: class DQN(nn.Module): #base model
    def __init__(self, num_inputs, num_actions, HIDDEN_LAYER_WIDTH):
        super(DQN, self).__init__()

        self.action_dim = num_actions

        self.layers = nn.Sequential(
            nn.Linear(num_inputs, HIDDEN_LAYER_WIDTH),
```

```

        nn.ReLU(),
        nn.Linear(HIDDEN_LAYER_WIDTH, HIDDEN_LAYER_WIDTH),
        nn.ReLU(),
        nn.Linear(HIDDEN_LAYER_WIDTH, num_actions)
    )

    def forward(self, x):
        return self.layers(x)

    def act(self, state, epsilon):
        with torch.no_grad():
            if random.random() > epsilon:
                state = torch.tensor(state, dtype=torch.float32).
→unsqueeze(dim=0).to(device)
                q_values = self.forward(state)
                action = q_values.max(dim=1)[1].item()
            else:
                action = random.randrange(self.action_dim)
        return action

```

```
[ ]: # e-greedy exploration
```

```

epsilon_start = config['EPSILON_START']
epsilon_final = config['EPSILON_FINAL']
epsilon_decay = config['EPSILON_DECAY']

epsilon_by_frame = lambda frame_idx: epsilon_final + (epsilon_start -
→epsilon_final) * math.exp(-1. * frame_idx / epsilon_decay)

```

```
[ ]: plt.plot([epsilon_by_frame(i) for i in range(10000)])
```

```
[ ]: # MODEL
```

```

if (config['MODEL_NAME']=='D1QN'):
    # only one NN for estimating Q-values
    model = DQN(env.observation_space.shape[0],
                env.action_space.n,
                config['HIDDEN_LAYER_WIDTH'])
    model = model.to(device)

elif (config['MODEL_NAME']=='DQN' or config['MODEL_NAME']=='D2QN'):
    # one inference model and one target model
    model = DQN(env.observation_space.shape[0],
                env.action_space.n,
                config['HIDDEN_LAYER_WIDTH'])
    model = model.to(device)

    target = DQN(env.observation_space.shape[0],

```

```

        env.action_space.n,
        config['HIDDEN_LAYER_WIDTH'])
    target = target.to(device)

else: #default model is D1QN
    # only one NN for estimating Q-values
    model = DQN(env.observation_space.shape[0],
                env.action_space.n,
                config['HIDDEN_LAYER_WIDTH'])
    model = model.to(device)

print(model)
summary(model,
        input_size=(env.observation_space.shape[0],),
        batch_size=config['BATCH_SIZE'],
        device='cuda' if USE_CUDA else 'cpu' )

```

```

[:]: # OPTIMIZER
if (config['OPTIMIZER']=='Adam'):
    optimizer = optim.Adam(model.parameters(),
                           lr=config['LEARNING_RATE'])
elif (config['OPTIMIZER']=='SGD'):
    optimizer = optim.SGD(model.parameters(),
                          lr=config['LEARNING_RATE'])
else: #default optimizer is Adam
    optimizer = optim.Adam(model.parameters(),
                           lr=config['LEARNING_RATE'])

```

```

[:]: # CRITERION
if (config['CRITERION']=='MSE'):
    criterion = nn.MSELoss()
elif (config['CRITERION']=='HUBER'):
    criterion = nn.SmoothL1Loss()
else: #default criterion is MSELoss
    criterion = nn.MSELoss()

```

```

[:]: # REPLAY BUFFER
replay_buffer = ReplayBuffer(capacity=config['REPLAY_BUFFER_SIZE'])

```

```

[:]: def update_target(current_model, target_model):
    target.load_state_dict(model.state_dict())

def compute_td_loss(batch_size):
    state, action, reward, next_state, done = replay_buffer.sample(batch_size)

    state = torch.tensor(np.float32(state), dtype=torch.float32).
    →to(device)

```

```

    next_state = torch.tensor(np.float32(next_state) ,dtype=torch.float32,
→requires_grad=False).to(device)
    action      = torch.tensor(action                        ,dtype=torch.long).
→to(device)
    reward      = torch.tensor(reward                      ,dtype=torch.float32).
→to(device)
    done        = torch.tensor(done                        ,dtype=torch.float32).
→to(device)

    q_values = model(state)
    q_value  = q_values.gather(dim=1, index=action.unsqueeze(dim=1)).
→squeeze(dim=1)

    #next_q_value
    if (config['MODEL_NAME']=='D1QN'):
        next_q_values = model(next_state)
        next_q_value  = next_q_values.max(dim=1)[0]

    elif (config['MODEL_NAME']=='DQN'):
        next_q_values = target(next_state)
        next_q_value  = next_q_values.max(dim=1)[0]

    elif (config['MODEL_NAME']=='D2QN'):
        next_q_values = model(next_state) #all q-values from current model
        next_q_target_values = target(next_state) #all q-values from target
→model
        next_q_value = next_q_target_values.gather(dim=1,
                                                    index=torch.
→max(next_q_values, dim=1)[1].unsqueeze(dim=1)).squeeze(dim=1)
        #q-values from target model by acting greedily on current model (double
→dqn)

    else: #Default is D1QN
        next_q_values = model(next_state)
        next_q_value  = next_q_values.max(dim=1)[0]

    expected_q_value = reward + gamma * next_q_value * (1 - done)

    loss = criterion(q_value, expected_q_value)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    return loss.to('cpu')

```

```
[ ]: def plot(frame_idx, rewards, losses):
    clear_output(True)
    plt.figure(figsize=(20,5))
    plt.subplot(131)
    plt.title('frame %s. reward: %s' % (frame_idx, np.mean(rewards[-10:])))
    plt.plot(rewards)
    plt.subplot(132)
    plt.title('loss')
    plt.plot(losses)
    plt.show()
```

```
[ ]: if (config['MODEL_NAME']=='DQN' or config['MODEL_NAME']=='D2QN'):
    update_target(model, target)
```

```
[ ]: %%time

# Training
num_frames = config['TIMESTEPS']
batch_size = config['BATCH_SIZE']
gamma      = config['GAMMA']

losses = []
all_rewards = []
episode_reward = 0

state = env.reset()

for frame_idx in range(1, num_frames + 1):
    epsilon = epsilon_by_frame(frame_idx)
    action = model.act(state, epsilon)

    next_state, reward, done, _ = env.step(action)
    replay_buffer.push(state, action, reward, next_state, done)

    state = next_state
    episode_reward += reward

    if done:
        writer.add_scalar('episode_reward', episode_reward,
→global_step=frame_idx)
        state = env.reset()
        all_rewards.append(episode_reward)
        episode_reward = 0

    if len(replay_buffer) > batch_size:
        loss = compute_td_loss(batch_size)
        losses.append(loss.item())
        writer.add_scalar('loss', loss.item(), global_step=frame_idx)
```

```

        for name, param in model.named_parameters():
            if param.requires_grad:
                writer.add_histogram('model_' + name, param.data,
→global_step=frame_idx)

        if (config['MODEL_NAME']=='DQN' or config['MODEL_NAME']=='D2QN'):
            for name, param in target.named_parameters():
                if param.requires_grad:
                    writer.add_histogram('target_' + name, param.data,
→global_step=frame_idx)

        if (config['MODEL_NAME']=='DQN' or config['MODEL_NAME']=='D2QN'):
            if frame_idx % config['TARGET_UPDATE_FREQ'] == 0:
                update_target(model, target)

writer.close()

```

```
[ ]: !jupyter nbconvert --to pdf dqn-cartpole.ipynb
```